

Heterogeneous Bipartite Graph Neural Networks for Lattice Quantum Field Theory

Joshua Paul Brehm^{1,*}

¹*Independent Researcher*

(Dated: April 2, 2026)

We introduce a heterogeneous bipartite graph neural network architecture that separates *space-time geometry* from *field content* into distinct node types connected by typed edges—mirroring the geometry–matter separation fundamental to continuum quantum field theory. Unlike existing approaches that embed all information on spacetime nodes in a homogeneous graph, our architecture represents fields as a separate layer coupled to the lattice through typed “inhabits” edges and three-stage message passing. We benchmark on two-dimensional scalar ϕ^4 theory in the Ising universality class ($\nu = 1$), comparing against homogeneous GNN, convolutional, and MLP baselines. We show that homogeneous GNN representations suffer a structural failure mode—message passing destroys the field signal without ad hoc skip connections—which the bipartite design eliminates by construction. The model achieves near-perfect action prediction ($r > 0.999$) across lattice sizes from 8×8 to 64×64 and reproduces finite-size scaling behavior including susceptibility peak growth, order parameter S-curves, and correlation length crossing consistent with the exact critical exponent. The bipartite structure naturally accommodates multiple field types (scalar, gauge, fermion) at the same lattice site, providing a direct path to gauge theories and lattice QCD without architectural changes.

I. INTRODUCTION

Lattice quantum field theory is the primary non-perturbative tool for studying strongly coupled quantum systems, underpinning our quantitative understanding of hadron physics through lattice QCD [1]. Despite decades of algorithmic advances, lattice calculations face persistent computational bottlenecks: critical slowing down near phase transitions, the fermion sign problem at finite baryon density, and the inability to directly simulate real-time dynamics in Minkowski spacetime.

Machine learning methods have shown promise in addressing several of these challenges. Normalizing flows can accelerate sampling by learning the Boltzmann distribution [2, 3], while graph neural networks (GNNs) can learn lattice observables directly from field configurations [4, 5]. However, existing GNN approaches treat the lattice as a *homogeneous graph*, embedding field values (scalar, gauge links, spinor components) as features on spacetime nodes. This conflates two fundamentally different objects: the geometry of spacetime and the dynamical content of quantum fields.

In continuum QFT, the action functional separates geometry from matter:

$$S[\phi, g] = \int d^d x \sqrt{g} \left[\frac{1}{2} g^{\mu\nu} \partial_\mu \phi \partial_\nu \phi + V(\phi) \right], \quad (1)$$

where the metric $g_{\mu\nu}$ encodes spacetime geometry and ϕ encodes field content. These couple through the action but are distinct mathematical objects living in different spaces.

We propose a *heterogeneous bipartite graph* architecture that mirrors this separation on the lattice. Space-

time nodes encode lattice geometry (coordinates, spacing, boundary conditions), while field nodes encode dynamical degrees of freedom (scalar values, gauge link matrices, spinor components). The two layers communicate through typed “inhabits” edges and three-stage message passing. This design has several advantages:

1. **Natural multi-field handling:** Multiple field types at the same site are distinct nodes with separate encoders, avoiding feature concatenation.
2. **Dynamic geometry:** Spacetime node positions can be made learnable, enabling adaptive mesh refinement.
3. **Extensibility:** Adding gauge or fermion fields requires only new encoder modules, not architectural changes.
4. **Physical interpretability:** The bipartite structure preserves the geometry–matter distinction through the entire network.

We benchmark this architecture on two-dimensional scalar ϕ^4 theory, the simplest interacting field theory with a well-understood phase transition in the Ising universality class ($\nu = 1$ exactly). This provides a rigorous test: the model must learn the action functional, reproduce the phase transition, and yield correct finite-size scaling across lattice sizes $L = 16, 32, 64$.

A. Related Work and Novelty

Machine learning approaches to lattice field theory have developed along several directions. Normalizing flows learn the Boltzmann distribution for efficient sampling [2, 3, 6]. Gauge-equivariant convolutional neural

* joshua.paul.brehm@gmail.com

networks enforce exact lattice symmetries in the network architecture [4, 7]. Graph neural networks have been applied to learn field-theoretic observables [5], and relational inductive biases have proven valuable in broader physics contexts [8, 9].

These approaches share a common representation: field variables are attached directly to lattice sites or links, with all information embedded on a single node type in a homogeneous graph (or equivalently on a regular grid). While effective, this conflates spacetime geometry with field content—objects that are mathematically distinct in continuum QFT.

Our contribution is a *heterogeneous bipartite graph formulation* in which spacetime geometry and field degrees of freedom are represented as distinct node types connected by typed edges. This mirrors the fiber bundle structure of continuum field theory: spacetime nodes play the role of the discrete base manifold M , field nodes correspond to sections of the associated fiber bundle, and the “inhabits” edges encode the projection map $\pi : E \rightarrow M$ associating each field value to its spacetime site. The three-stage message-passing protocol (field $\rightarrow$ spacetime, spacetime $\rightarrow$ spacetime, spacetime $\rightarrow$ field) acts as a discrete analog of the field–geometry coupling encoded in the action functional. The resulting architecture naturally supports multiple field types at a single lattice site without feature concatenation and provides a modular pathway to more complex theories, including gauge fields and fermions.

II. ARCHITECTURE

A. Heterogeneous Bipartite Graph

For a scalar field theory on an $L \times L$ lattice with periodic boundary conditions, the graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ consists of two node types and three edge types:

Node types:

- *Spacetime nodes* ($N = L^2$): features are coordinates (x_1, x_2) .
- *Scalar field nodes* ($N = L^2$): features are the field value (ϕ_i) at each site.

Edge types:

- *Adjacent* (spacetime $\rightarrow$ spacetime): $2dN$ directed edges encoding the lattice connectivity, with displacement vectors $a\hat{e}_\mu$ as features. These encode both the direction and physical distance between neighboring sites, placing metric information on the edges where it naturally belongs.
- *Inhabits* (field $\rightarrow$ spacetime): N edges connecting each field node to its corresponding spacetime site.
- *Inhabits-inverse* (spacetime $\rightarrow$ field): reverse of the above, enabling spacetime-to-field message flow.




FIG. 1. 3D view of the heterogeneous bipartite graph for a 4×4 lattice. Blue spheres (upper plane): spacetime nodes with coordinate features (x_1, x_2) . Orange spheres (lower plane): scalar field nodes with ϕ values. Solid blue lines: adjacency edges encoding lattice connectivity. Dashed orange lines: inhabits edges coupling each field node to its spacetime site. Reverse inhabits edges (not shown) enable bidirectional message passing.

For an $L \times L$ lattice, the graph has $2L^2$ nodes and $4L^2 + 2L^2 = 6L^2$ directed edges (for $d = 2$). The graph is implemented using PyTorch Geometric’s `HeteroData` container.

B. Model: HeteroGNN

The model consists of three stages: type-specific encoding, iterative three-stage message passing, and a physics-informed readout head.

Encoding. Each node and edge type has a dedicated MLP encoder that projects raw features to a shared hidden dimension H :

$$\mathbf{h}_i^{(\text{st})} = \text{MLP}_{\text{st}}(x_i^1, x_i^2) \in \mathbb{R}^H, \quad (2)$$

$$\mathbf{h}_i^{(\text{field})} = \text{MLP}_{\text{field}}(\phi_i) \in \mathbb{R}^H, \quad (3)$$

$$\mathbf{e}_{ij}^{(\text{adj})} = \text{MLP}_{\text{edge}}(a\hat{e}_\mu) \in \mathbb{R}^H. \quad (4)$$

Three-stage message passing. Each message-passing block performs three sequential steps with residual connections and layer normalization:

Stage 1: Field $\rightarrow$ Spacetime. Field information is aggregated at each spacetime site via the inhabits edges:

$$\mathbf{h}_i^{(\text{st})} += \text{MLP}_1 \left(\mathbf{h}_i^{(\text{st})} \parallel \sum_{j \in \mathcal{N}_{\text{inh}}(i)} \mathbf{h}_j^{(\text{field})} \right). \quad (5)$$

Stage 2: Spacetime $\rightarrow$ Spacetime. Information propagates along the lattice via adjacency edges, analogous to discretized spatial derivatives:

$$\mathbf{h}_i^{(\text{st})} += \text{MLP}_2 \left(\mathbf{h}_i^{(\text{st})} \parallel \sum_{j \in \mathcal{N}_{\text{adj}}(i)} [\mathbf{h}_j^{(\text{st})} \parallel \mathbf{e}_{ij}] \right). \quad (6)$$

Stage 3: Spacetime $\rightarrow$ Field. Updated geometry information flows back to field nodes:

$$\mathbf{h}_i^{(\text{field})} += \text{MLP}_3 \left(\mathbf{h}_i^{(\text{field})} \parallel \sum_{j \in \mathcal{N}_{\text{inh}}^{-1}(i)} \mathbf{h}_j^{(\text{st})} \right). \quad (7)$$

Each step includes residual connections ($+=$) and layer normalization to mitigate over-smoothing. We stack B such blocks; with B blocks, the receptive field spans B lattice spacings.

Readout: Action head. The total Euclidean action is predicted as a sum of local action densities:

$$\hat{S}_E[\phi] = \sum_{i=1}^N \text{MLP}_{\text{action}} \left(\mathbf{h}_i^{(\text{st})} \parallel \mathbf{h}_i^{(\text{field})} \right) \cdot a^d, \quad (8)$$

where the a^d factor provides the lattice volume element, converting the discrete sum into the continuum integral $\int d^d x \mathcal{L}$. This factor is not merely cosmetic: ablation experiments (Sec. V A) show that removing it causes catastrophic failure at non-unit lattice spacings, as the MLP cannot efficiently learn to compensate for a scale factor spanning several orders of magnitude. The per-site decomposition mirrors the structure of the lattice action.

III. BENCHMARK: SCALAR ϕ^4 THEORY

A. Theory and Observables

We study the Euclidean action on a two-dimensional square lattice:

$$S_E[\phi] = \sum_x a^2 \left[\frac{1}{2} \sum_{\mu} \frac{(\phi_{x+\hat{\mu}} - \phi_x)^2}{a^2} + \frac{m^2}{2} \phi_x^2 + \lambda \phi_x^4 \right], \quad (9)$$

with periodic boundary conditions. For $\lambda = 0.5$, this theory undergoes a second-order phase transition in the Ising universality class as m^2 is varied, with critical exponent $\nu = 1$ (exact, from conformal field theory [10]).

Key observables for the finite-size scaling analysis are:

- **Order parameter:** $|M| = |\langle \phi \rangle| = |V^{-1} \sum_x \phi_x|$
- **Susceptibility:** $\chi = V(\langle M^2 \rangle - \langle |M| \rangle^2)$
- **Correlation length** (second-moment estimator):

$$\xi = \frac{1}{2 \sin(\pi/L)} \sqrt{\frac{\tilde{G}(0)}{\tilde{G}(k_{\min})} - 1}, \quad (10)$$

where $\tilde{G}(k)$ is the Fourier-transformed connected correlator and $k_{\min} = 2\pi/L$.

We implement the correlation length estimator using the full two-dimensional FFT of the field configurations:

$$\tilde{G}(0) = \chi = V \text{Var}(M), \quad (11)$$

$$\tilde{G}(k_{\min}) = \langle |\tilde{\phi}(k_{\min}, 0)|^2 \rangle / V, \quad (12)$$

which uses all spatial correlations and avoids the systematic underestimation of one-dimensional angle-averaged approaches.

B. Monte Carlo Data Generation

Field configurations are generated using the Metropolis–Hastings algorithm with single-site Gaussian proposals. For lattices of size $L \geq 32$, we employ a *checkerboard decomposition*: sites are partitioned into two sublattices by coordinate parity, and all sites on the same sublattice are updated simultaneously in a single vectorized operation. This preserves detailed balance (no two updated sites are neighbors) while achieving 10–50 $\times$ speedup over sequential updates.

We generate 5000 configurations at each lattice size ($L = 8, 16, 32, 64$) for training, with $m^2 = -0.5$ and $\lambda = 0.5$. Configurations are separated by $n_{\text{sweeps}} = 10$ full sweeps (scaled to 40 for $L = 64$) to reduce autocorrelation, following 1000–1500 thermalization sweeps. Acceptance rates are maintained near the optimal $\sim 47\%$ by tuning the proposal step size.

C. Training

The GNN is trained to minimize the mean squared error between predicted and exact Euclidean actions:

$$\mathcal{L} = \frac{1}{N_{\text{batch}}} \sum_{i=1}^{N_{\text{batch}}} \left(\hat{S}_E[\phi_i] - S_E[\phi_i] \right)^2. \quad (13)$$

Model hyperparameters: hidden dimension $H = 64$, $B = 3$ message-passing blocks, GELU activation, 2-layer MLPs in encoders. Training uses AdamW optimization with cosine annealing over 150 epochs, batch size 32, and an 80/20 train/validation split.

IV. RESULTS

A. Action Prediction

Table I summarizes the action prediction results. The model achieves near-perfect correlation on both lattice sizes, demonstrating that the heterogeneous bipartite architecture can faithfully represent the action functional of scalar ϕ^4 theory. The slight decrease in accuracy at

TABLE I. Action prediction accuracy across lattice sizes. Pearson correlation r and mean relative error are computed on held-out validation sets.

Lattice	Sites	Pearson r	Relative Error
16×16	256	1.0000	$< 0.1\%$
64×64	4096	0.9994	$< 0.5\%$

TABLE II. Architecture comparison on 16×16 lattice ($m^2 = -0.5$, $\lambda = 0.5$). All models trained for 150 epochs on the same data split.

Model	Params	Pearson r	Rel. Error
HeteroGNN (ours)	150,337	0.99998	0.04%
Homogeneous GNN	67,073	0.99985	0.12%
Lattice CNN	29,153	0.99855	0.39%
MLP	197,633	0.346	7.4%

64×64 is expected: the graph has $16\times$ more nodes, and the receptive field of $B = 3$ message-passing blocks covers a smaller fraction of the lattice.

B. Architecture Comparison

To isolate the contribution of the bipartite architecture, we compare against three baselines trained on the same 16×16 data with matched training protocols (150 epochs, AdamW, cosine annealing):

- **Homogeneous GNN:** Field values concatenated with spacetime coordinates on a single node type; same message-passing depth ($B = 3$) and hidden dimension ($H = 64$), with a skip connection from the initial encoding to the readout to prevent over-smoothing of the field signal.
- **Lattice CNN:** Periodic-padded 2D convolutions treating $\phi(x)$ as a single-channel image (4 layers, 32 channels).
- **MLP:** Flattened field configuration fed through a 3-layer feedforward network ($H = 256$).

Table II shows the results. The most significant finding is a *structural failure mode* in the homogeneous GNN: without an explicit skip connection from its initial encoding to the readout head, message passing destroys the field signal entirely ($r < 0.05$). This workaround is not physically motivated—it is an engineering patch required because field values and geometry features share a single node representation that is progressively smoothed during message passing. The bipartite architecture avoids this failure mode by construction: field nodes maintain a separate representation throughout the network, ensuring the field signal reaches the readout regardless of message-passing depth. Note that the per-stage residual

TABLE III. Generalization across coupling values. The model was trained at $m^2 = -0.5$ (marked $*$) and evaluated on configurations at other m^2 values without retraining.

m^2	Pearson r	Rel. Error
-0.3	0.9993	6.1%
-0.5*	0.9999	0.04%
-0.7	0.9990	7.1%
-1.0	0.9915	20.6%
-1.5	0.9324	58.0%
-2.0	0.5846	211.5%

connections in the HeteroGNN (Sec. II B) serve a different purpose: they stabilize gradient flow within each message-passing block, whereas the homogeneous skip connection bypasses message passing entirely to rescue a signal that would otherwise be destroyed.

Beyond architectural stability, all spatially structured models achieve $r > 0.998$, confirming that the ϕ^4 action is learnable as a local functional. The HeteroGNN achieves the lowest relative error (0.04%), a $3\times$ improvement over the homogeneous GNN with its skip connection (0.12%). The MLP baseline fails without any spatial structure ($r = 0.35$), confirming that the lattice topology is essential.

C. Coupling Generalization

To test whether the model has learned a transferable action functional rather than memorizing a single action landscape, we evaluate the model trained at $m^2 = -0.5$ on configurations generated at different m^2 values (same $\lambda = 0.5$).

Table III shows that the model generalizes well to nearby couplings: for $|m^2 - m_{\text{train}}^2| \leq 0.5$, the Pearson correlation remains above 0.999, indicating the learned action functional transfers across this range. Performance degrades gracefully for more distant couplings, with $r = 0.99$ at $m^2 = -1.0$ and $r = 0.93$ at $m^2 = -1.5$. The sharp drop at $m^2 = -2.0$ ($r = 0.58$) coincides with the approach to the critical region ($m_c^2 \approx -2.45$), where the field distribution changes qualitatively and extrapolation from the symmetric phase becomes unreliable. Training on a mixture of couplings would likely extend the generalization range substantially.

D. Phase Transition and Finite-Size Scaling

We perform a coupling sweep over $m^2 \in [-2.8, -1.5]$ at $\lambda = 0.5$ for $L = 16, 32, 64$ using 2000 configurations per point with warm-starting between successive m^2 values.

Figure 2 shows the three canonical finite-size scaling observables. The order parameter exhibits the expected S-curve behavior, with the transition sharpening progressively from $L = 16$ to $L = 64$. The susceptibility peaks




FIG. 2. Finite-size scaling results for $\lambda = 0.5$. **Left:** Order parameter $|\phi|$ showing progressive sharpening of the phase transition with increasing L . **Center:** Susceptibility χ with peaks growing as $\chi_{\max} \sim L^{\gamma/\nu}$ and converging toward m_c^2 . **Right:** Correlation length ratio ξ/L with crossing of the $L = 16$ and $L = 32$ curves near $m_c^2 \approx -2.35$.

grow with volume— $\chi_{\max} \approx 10, 25, 125$ for $L = 16, 32, 64$ respectively—consistent with the scaling $\chi_{\max} \sim L^{\gamma/\nu}$ with $\gamma/\nu \approx 1.75$ (the exact two-dimensional Ising value).

The ξ/L ratio shows a crossing of the $L = 16$ and $L = 32$ curves near $m_c^2 \approx -2.35$, with the $L = 32$ and $L = 64$ curves approaching intersection at $m^2 \approx -2.5$. The remaining gap at $L = 64$ is attributable to critical slowing down of the local Metropolis algorithm: autocorrelation times grow as $\tau \sim \xi^z$ with dynamic exponent $z \approx 2$, degrading the effective statistics near the critical point for the largest lattice.

E. Critical Exponent Extraction

From the ξ/L crossing and data collapse analysis, we obtain:

$$\nu = 1.1 \pm 0.8 \quad (\text{exact: } \nu = 1), \quad (14)$$

$$m_c^2 = -2.45 \pm 0.10 \quad (\lambda = 0.5). \quad (15)$$

We emphasize that this serves as a *consistency check* rather than a precision measurement: the central value is compatible with the exact Ising result $\nu = 1$, but the large uncertainty precludes a strong quantitative claim. The dominant source of error is not the GNN architecture but the Monte Carlo sampling itself. The local Metropolis algorithm suffers from critical slowing down ($\tau \sim \xi^z$, $z \approx 2$), which degrades effective statistics precisely where they matter most—near the critical point. A cluster algorithm (Wolff [11] or Swendsen–Wang [12]) would reduce z to near zero and dramatically tighten the ν estimate, but is beyond the scope of this initial architecture demonstration. Our purpose here is to show that the bipartite GNN *captures* the physics of the phase transition, not to compete with dedicated finite-size scaling studies.

V. DISCUSSION

A. Advantages of the Bipartite Architecture

The heterogeneous bipartite graph offers several structural advantages over homogeneous approaches:

Multi-field extensibility. In the homogeneous approach, adding a gauge field $U_\mu(x) \in \text{SU}(N)$ requires concatenating its matrix elements with the scalar field value at each node, breaking the natural group structure. In our architecture, gauge fields become a separate node type with their own encoder, connected to space-time via typed edges. The three-stage message-passing blocks handle arbitrary numbers of field types without modification.

Physical interpretability. The action readout (Eq. 8) decomposes the total action into per-site contributions $\hat{s}(x)$ that can be directly compared with the lattice action density, enabling local error analysis.

Metric information placement. The lattice spacing a is a metric quantity—it defines the physical distance between neighboring sites. We encode it on adjacency edges as displacement vectors $a\hat{e}_\mu$ rather than as a node feature, reflecting the fact that distances are properties of connections, not of individual points. To validate this choice, we tested four architecture variants (lattice spacing in nodes vs. edges, with and without the explicit a^d volume factor in Eq. 8) across lattice sizes 8×8 through 64×64 at varying lattice spacings $a = 1/L$. Two findings emerged. First, the a^d volume scaling is essential: without it, performance collapses as a decreases (Pearson r drops from > 0.999 to -0.26 at 64×64 with $a = 1/64$), because the MLP readout cannot learn to compensate for a scale factor spanning four orders of magnitude. Second, the choice of whether a resides on nodes or edges has negligible effect on accuracy ($r > 0.999$ for both at all sizes), confirming that the placement is a design choice rather than a performance constraint. We choose edges because the metric naturally lives on links, and this design extends cleanly to gauge theories where link variables $U_\mu(x)$ share the same edge representation.

Over-smoothing mitigation. By separating field and spacetime representations, the architecture avoids mixing all information into a single node type during message passing. Residual connections and layer normalization further stabilize deep stacking.

Interpretation as a learned effective action. The action readout (Eq. 8) decomposes the total action into per-site contributions, and the finite receptive field of B message-passing blocks limits each contribution to information within Ba lattice spacings. The model therefore learns a *local effective action* operator—analogous to a Wilsonian effective action in which modes above a momentum cutoff $\Lambda \sim 1/(Ba)$ have been integrated out. In this view, the message-passing depth B plays the role of an implicit renormalization scale, controlling the resolution at which the theory is represented. This interpretation explains both the model’s success on the local ϕ^4

action and its degradation near criticality, where the relevant physics involves correlations at scales $\xi \gg Ba$ that lie below the effective cutoff.

B. Limitations

Critical slowing down. The Metropolis sampler’s autocorrelation time diverges at the critical point, limiting our ability to precisely determine ν . This is a sampling limitation, not an architectural one.

Receptive field and locality scale. With $B = 3$ blocks, the model’s receptive field spans $\ell = Ba = 3a$, defining a finite *locality scale* for the learned operator. This is sufficient for the ϕ^4 action, which is a strictly local functional of the field and its nearest-neighbor differences. However, near the critical point where $\xi \gg \ell$, the model cannot capture the long-range correlations that dominate critical behavior. The observed performance degradation at criticality (Sec. IV C) is therefore not merely a limitation but *expected physical behavior*: it confirms that the network has learned a genuinely local operator rather than memorizing global patterns. Increasing B enlarges the locality scale but risks over-smoothing, where distinct node representations collapse toward a common value. For theories requiring larger locality scales, possible mitigations include multi-scale pooling layers, skip connections across non-adjacent blocks, or attention mechanisms that allow selective long-range communication.

Lattice symmetries. The current architecture does not explicitly enforce discrete lattice symmetries (translations, C_4 rotations, reflections). Instead, these symmetries are learned implicitly from the data, which is generated with periodic boundary conditions that respect the full symmetry group. This is a deliberate design choice, not a limitation: gauge-equivariant architectures [4, 7] achieve exact symmetry by constraining the network weights, but this rigidity restricts them to regular lattices with specific boundary conditions. Our approach prioritizes *architectural flexibility*—the same model structure handles irregular meshes, non-periodic boundaries, and adaptive lattice refinements without modification. For the multi-theory extensibility that motivates this work (gauge fields, fermions, curved spacetimes), this flexibility is essential. Quantifying the learned equivariance error on regular lattices is left to future work.

VI. FUTURE DIRECTIONS

The bipartite architecture is designed as a foundation for progressively more complex theories:

Phase II: U(1) gauge theory with fermions. Gauge link variables $U_\mu(x) = e^{iA_\mu(x)} \in \text{U}(1)$ and Dirac spinors $\psi(x) \in \mathbb{C}^2$ are added as new field node types with dedicated encoders. The three-stage message passing naturally handles the tripartite field–gauge–

spacetime coupling. This phase will benchmark against the Schwinger model (QED in 1+1D), which has exact solutions for comparison.

Phase III: SU(3) Yang–Mills in 4D. The architecture extends to four-dimensional lattices by simply setting the lattice dimension parameter to 4. Gauge link encoders must respect the SU(3) group structure; we plan to use equivariant neural network layers for this purpose. The target is to reproduce established lattice QCD benchmarks (string tension, glueball spectrum).

Phase IV: Wick rotation bridge. Using the GNN’s learned representation of the Euclidean theory, we aim to construct a bridge to Minkowski spacetime via spectral function reconstruction, connecting imaginary-time lattice data to real-time physics.

Phase V: Native Minkowski dynamics. The ultimate goal is real-time simulation of quantum field dynamics in Minkowski spacetime, potentially using complex Langevin methods guided by the GNN’s learned action functional.

VII. CONCLUSION

We have introduced a heterogeneous bipartite graph neural network that separates spacetime geometry from field content—a representation inspired by the geometry–matter distinction fundamental to continuum QFT, and absent from existing lattice QFT architectures. On two-dimensional scalar ϕ^4 theory, the model achieves near-perfect action prediction ($r > 0.999$), reproduces the expected Ising universality class finite-size scaling, and yields critical exponent estimates consistent with exact values. Comparison against homogeneous GNN, CNN, and MLP baselines demonstrates the viability of the bipartite representation for lattice field theory.

The architecture’s key contribution is not improved accuracy on this benchmark but its *extensibility and physical alignment*: the same graph structure and message-passing protocol accommodate scalar, gauge, and fermion fields through modular encoder additions, while preserving the geometry–field distinction throughout the computation. Metric information (lattice spacing) is encoded on adjacency edges as displacement vectors, placing it where distances physically reside and enabling clean integration with gauge link variables in future phases. This provides a unified framework for progressing from toy models to physically relevant theories, with lattice QCD as the ultimate target.

Code and data are available at <https://github.com/crosis19/qft-graph>.

ACKNOWLEDGMENTS

Substantial portions of the code, analysis, and manuscript were developed in collaboration with Claude (Anthropic), an AI assistant. Pre-submission feedback

was obtained from ChatGPT (OpenAI) and Gemini (Google), which informed revisions to the experimental design and presentation.

-
- [1] K. G. Wilson, Confinement of quarks, *Physical Review D* **10**, 2445 (1974).
 - [2] G. Kanwar, M. S. Albergo, D. Boyda, K. Cranmer, D. C. Hackett, S. Racanière, D. J. Rezende, and P. E. Shanahan, Equivariant flow-based sampling for lattice gauge theory, *Physical Review Letters* **125**, 121601 (2020).
 - [3] D. Boyda, G. Kanwar, S. Racanière, D. J. Rezende, M. S. Albergo, K. Cranmer, D. C. Hackett, and P. E. Shanahan, Sampling using $SU(N)$ gauge equivariant flows, *Physical Review D* **103**, 074504 (2021).
 - [4] M. Favoni, A. Ipp, D. I. Müller, and D. Schuh, Lattice gauge equivariant convolutional neural networks, *Physical Review Letters* **128**, 032003 (2022).
 - [5] D. Bachtis, G. Aarts, and B. Lucini, Quantum field-theoretic machine learning, *Physical Review D* **104**, 075021 (2021).
 - [6] M. S. Albergo, G. Kanwar, and P. E. Shanahan, Flow-based generative models for Markov chain Monte Carlo in lattice field theory, *Physical Review D* **100**, 034515 (2019).
 - [7] D. Luo, S. Chen, G. Carleo, and B. K. Clark, Gauge equivariant neural networks for quantum lattice gauge theories, *Physical Review Letters* **127**, 276402 (2021).
 - [8] M. Cranmer, A. Sanchez-Gonzalez, P. Battaglia, R. Xu, K. Cranmer, D. Spergel, and S. Ho, Discovering symbolic and physical laws using graph networks, *arXiv preprint arXiv:2003.04630* (2020).
 - [9] P. W. Battaglia, J. B. Hamrick, V. Bapst, A. Sanchez-Gonzalez, V. Zambaldi, *et al.*, Relational inductive biases, deep learning, and graph networks, in *arXiv preprint arXiv:1806.01261* (2018).
 - [10] P. Di Francesco, P. Mathieu, and D. Sénéchal, *Conformal Field Theory* (Springer, 1997).
 - [11] U. Wolff, Collective Monte Carlo updating for spin systems, *Physical Review Letters* **62**, 361 (1989).
 - [12] R. H. Swendsen and J.-S. Wang, Nonuniversal critical dynamics in Monte Carlo simulations, *Physical Review Letters* **58**, 86 (1987).