

Heterogeneous Bipartite Graph Neural Networks for Lattice Quantum Field Theory

Joshua Paul Brehm^{1,*}

¹*Independent Researcher*

(Dated: July 9, 2026)

We introduce a heterogeneous bipartite graph neural network architecture that separates *space-time geometry* from *field content* into distinct node types connected by typed edges—mirroring the geometry–matter separation fundamental to continuum quantum field theory. Unlike existing approaches that embed all information on spacetime nodes in a homogeneous graph, our architecture represents fields as a separate layer coupled to the lattice through typed “inhabits” edges and three-stage message passing. We benchmark on two-dimensional scalar ϕ^4 theory in the Ising universality class ($\nu = 1$), comparing against homogeneous GNN, convolutional, and MLP baselines. We show that homogeneous GNN representations suffer a structural failure mode—message passing destroys the field signal without ad hoc skip connections—which the bipartite design eliminates by construction. The model achieves near-perfect action prediction across lattice sizes ($r > 0.9999$ at 8×8 and 16×16 ; median $r = 0.9992$ at 64×64), and a single model trained at 16×16 transfers to lattices from 8×8 to 64×64 with $r = 1.0000$ and no retraining. The underlying Monte Carlo pipeline reproduces Ising-class finite-size scaling—susceptibility peak growth with $\gamma/\nu = 1.57(12)$, order-parameter S-curves, and correlation-length crossings giving $\nu = 1.10(43)$, consistent with the exact $\nu = 1$ —validating the training distributions. The bipartite structure naturally accommodates multiple field types (scalar, gauge, fermion) at the same lattice site, providing a direct path to gauge theories and lattice QCD without architectural changes.

I. INTRODUCTION

Lattice quantum field theory is the primary non-perturbative tool for studying strongly coupled quantum systems, underpinning our quantitative understanding of hadron physics through lattice QCD [1]. Despite decades of algorithmic advances, lattice calculations face persistent computational bottlenecks: critical slowing down near phase transitions, the fermion sign problem at finite baryon density, and the inability to directly simulate real-time dynamics in Minkowski spacetime.

Machine learning methods have shown promise in addressing several of these challenges. Normalizing flows can accelerate sampling by learning the Boltzmann distribution [2, 3], while graph neural networks (GNNs) can learn lattice observables directly from field configurations [4, 5]. However, existing GNN approaches treat the lattice as a *homogeneous graph*, embedding field values (scalar, gauge links, spinor components) as features on spacetime nodes. This conflates two fundamentally different objects: the geometry of spacetime and the dynamical content of quantum fields.

In continuum QFT, the action functional separates geometry from matter:

$$S[\phi, g] = \int d^d x \sqrt{g} \left[\frac{1}{2} g^{\mu\nu} \partial_\mu \phi \partial_\nu \phi + V(\phi) \right], \quad (1)$$

where the metric $g_{\mu\nu}$ encodes spacetime geometry and ϕ encodes field content. These couple through the action

but are distinct mathematical objects living in different spaces.

We propose a *heterogeneous bipartite graph* architecture that mirrors this separation on the lattice. Spacetime nodes encode lattice geometry (coordinates, spacing, boundary conditions), while field nodes encode dynamical degrees of freedom (scalar values, gauge link matrices, spinor components). The two layers communicate through typed “inhabits” edges and three-stage message passing. This design has several advantages:

1. **Natural multi-field handling:** Multiple field types at the same site are distinct nodes with separate encoders, avoiding feature concatenation.
2. **Dynamic geometry:** Spacetime node positions can be made learnable, enabling adaptive mesh refinement.
3. **Extensibility:** Adding gauge or fermion fields requires only new encoder modules, not architectural changes.
4. **Physical interpretability:** The bipartite structure preserves the geometry–matter distinction through the entire network.

We benchmark this architecture on two-dimensional scalar ϕ^4 theory, the simplest interacting field theory with a well-understood phase transition in the Ising universality class ($\nu = 1$ exactly). This provides a rigorous test: the model must learn the action functional, reproduce the phase transition, and yield correct finite-size scaling across lattice sizes $L = 16, 32, 64$.

* joshua.paul.brehm@gmail.com

A. Related Work and Novelty

Machine learning approaches to lattice field theory have developed along several directions (see Ref. [6] for a recent review). Normalizing flows learn the Boltzmann distribution for efficient sampling [2, 3, 7]. Gauge-equivariant and gauge-invariant neural networks enforce exact lattice symmetries in the network architecture [4, 8, 9]. Graph neural networks have been applied to learn field-theoretic observables [5], and relational inductive biases have proven valuable in broader physics contexts [10, 11]. In the machine-learning literature, heterogeneous graphs with typed nodes and relations are a well-developed modeling tool [12, 13], but to our knowledge they have not been applied to the geometry-field decomposition of lattice field theory.

These approaches share a common representation: field variables are attached directly to lattice sites or links, with all information embedded on a single node type in a homogeneous graph (or equivalently on a regular grid). While effective, this conflates spacetime geometry with field content—objects that are mathematically distinct in continuum QFT.

Our contribution is a *heterogeneous bipartite graph formulation* in which spacetime geometry and field degrees of freedom are represented as distinct node types connected by typed edges. This structure is loosely analogous to a fiber bundle: spacetime nodes play the role of the discrete base manifold M , each field node represents a point in the fiber over its site (so that a field configuration—the collection of all field-node values—is a discretized section), and the “inhabits” edges encode the projection $\pi : E \rightarrow M$ associating each field value to its spacetime site. The three-stage message-passing protocol (field $\rightarrow$ spacetime, spacetime $\rightarrow$ spacetime, spacetime $\rightarrow$ field) acts as a discrete analog of the field-geometry coupling encoded in the action functional. The resulting architecture naturally supports multiple field types at a single lattice site without feature concatenation and provides a modular pathway to more complex theories, including gauge fields and fermions.

II. ARCHITECTURE

A. Heterogeneous Bipartite Graph

For a scalar field theory on an $L \times L$ lattice with periodic boundary conditions, the graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ consists of two node types and three edge types:

Node types:

- *Spacetime nodes* ($N = L^2$): features are coordinates (x_1, x_2) .
- *Scalar field nodes* ($N = L^2$): features are the field value (ϕ_i) at each site.

Edge types:




FIG. 1. 3D view of the heterogeneous bipartite graph for a 4×4 lattice. Blue spheres (upper plane): spacetime nodes with coordinate features (x_1, x_2) . Orange spheres (lower plane): scalar field nodes with ϕ values. Solid blue lines: adjacency edges encoding lattice connectivity. Dashed orange lines: inhabits edges coupling each field node to its spacetime site. Reverse inhabits edges (not shown) enable bidirectional message passing.

- *Adjacent* (spacetime $\rightarrow$ spacetime): $2dN$ directed edges encoding the lattice connectivity, with displacement vectors $a\hat{e}_\mu$ as features. These encode both the direction and physical distance between neighboring sites, placing metric information on the edges where it naturally belongs.
- *Inhabits* (field $\rightarrow$ spacetime): N edges connecting each field node to its corresponding spacetime site.
- *Inhabits-inverse* (spacetime $\rightarrow$ field): reverse of the above, enabling spacetime-to-field message flow.

For an $L \times L$ lattice, the graph has $2L^2$ nodes and $4L^2 + 2L^2 = 6L^2$ directed edges (for $d = 2$). The graph is implemented using PyTorch Geometric’s [14] `HeteroData` container.

B. Model: HeteroGNN

The model consists of three stages: type-specific encoding, iterative three-stage message passing, and a physics-informed readout head.

Encoding. Each node and edge type has a dedicated MLP encoder that projects raw features to a shared hid-

den dimension H :

$$\mathbf{h}_i^{(\text{st})} = \text{MLP}_{\text{st}}(x_i^1, x_i^2) \in \mathbb{R}^H, \quad (2)$$

$$\mathbf{h}_i^{(\text{field})} = \text{MLP}_{\text{field}}(\phi_i) \in \mathbb{R}^H, \quad (3)$$

$$\mathbf{e}_{ij}^{(\text{adj})} = \text{MLP}_{\text{edge}}(a\hat{e}_\mu) \in \mathbb{R}^H. \quad (4)$$

Three-stage message passing. Each message-passing block performs three sequential steps with residual connections and layer normalization:

Stage 1: Field $\rightarrow$ Spacetime. Field information is aggregated at each spacetime site via the inhabits edges:

$$\mathbf{h}_i^{(\text{st})} += \text{MLP}_1 \left(\mathbf{h}_i^{(\text{st})} \parallel \sum_{j \in \mathcal{N}_{\text{inh}}(i)} \mathbf{h}_j^{(\text{field})} \right). \quad (5)$$

Stage 2: Spacetime $\rightarrow$ Spacetime. Information propagates along the lattice via adjacency edges, analogous to discretized spatial derivatives:

$$\mathbf{h}_i^{(\text{st})} += \text{MLP}_2 \left(\mathbf{h}_i^{(\text{st})} \parallel \sum_{j \in \mathcal{N}_{\text{adj}}(i)} [\mathbf{h}_j^{(\text{st})} \parallel \mathbf{e}_{ij}] \right). \quad (6)$$

Stage 3: Spacetime $\rightarrow$ Field. Updated geometry information flows back to field nodes:

$$\mathbf{h}_i^{(\text{field})} += \text{MLP}_3 \left(\mathbf{h}_i^{(\text{field})} \parallel \sum_{j \in \mathcal{N}_{\text{inh}}^{-1}(i)} \mathbf{h}_j^{(\text{st})} \right). \quad (7)$$

Each step includes residual connections ($+=$) and layer normalization to mitigate over-smoothing. We stack B such blocks; with B blocks, the receptive field spans B lattice spacings.

Readout: Action head. The total Euclidean action is predicted as a sum of local action densities:

$$\hat{S}_E[\phi] = \sum_{i=1}^N \text{MLP}_{\text{action}}(\mathbf{h}_i^{(\text{st})} \parallel \mathbf{h}_i^{(\text{field})}) \cdot a^d, \quad (8)$$

where the a^d factor provides the lattice volume element, converting the discrete sum into the continuum integral $\int d^d x \mathcal{L}$. This factor is not merely cosmetic: ablation experiments (Sec. V A) show that removing it causes catastrophic failure at non-unit lattice spacings, as the MLP cannot efficiently learn to compensate for a scale factor spanning several orders of magnitude. The per-site decomposition mirrors the structure of the lattice action.

III. BENCHMARK: SCALAR ϕ^4 THEORY

A. Theory and Observables

We study the Euclidean action on a two-dimensional square lattice:

$$S_E[\phi] = \sum_x a^2 \left[\frac{1}{2} \sum_\mu \frac{(\phi_{x+\hat{\mu}} - \phi_x)^2}{a^2} + \frac{m^2}{2} \phi_x^2 + \lambda \phi_x^4 \right], \quad (9)$$

with periodic boundary conditions. For $\lambda = 0.5$, this theory undergoes a second-order phase transition in the Ising universality class as m^2 is varied, with critical exponent $\nu = 1$ (exact, from conformal field theory [15]).

Key observables for the finite-size scaling analysis are:

- **Order parameter:** $|M| = |\langle \phi \rangle| = |V^{-1} \sum_x \phi_x|$
- **Susceptibility:** $\chi = V(\langle M^2 \rangle - \langle |M| \rangle^2)$
- **Correlation length** (second-moment estimator):

$$\xi = \frac{1}{2 \sin(\pi/L)} \sqrt{\frac{\tilde{G}(0)}{\tilde{G}(k_{\min})} - 1}, \quad (10)$$

where $\tilde{G}(k)$ is the Fourier-transformed connected correlator and $k_{\min} = 2\pi/L$.

We implement the correlation length estimator using the full two-dimensional FFT of the field configurations:

$$\tilde{G}(0) = \chi = V(\langle M^2 \rangle - \langle |M| \rangle^2), \quad (11)$$

$$\tilde{G}(k_{\min}) = \langle |\tilde{\phi}(k_{\min}, 0)|^2 \rangle / V, \quad (12)$$

which uses all spatial correlations and avoids the systematic underestimation of one-dimensional angle-averaged approaches. The $k \neq 0$ modes need no subtraction since $\langle \phi(k \neq 0) \rangle = 0$, but the $k = 0$ mode is subtraction-dependent: the $\langle |M| \rangle$ -subtracted form degenerates in the symmetric phase (for Gaussian M , $\langle |M| \rangle^2 = (2/\pi) \langle M^2 \rangle$, driving the ratio in Eq. (10) negative), while the alternative $V \text{Var}(M)$ form is contaminated at finite volume by rare tunneling between the two ordered vacua in the broken phase. The ξ/L crossings and scaling collapse therefore use a two-momentum variant built entirely from $k \neq 0$ modes,

$$\xi^{-2} = \frac{\hat{k}_2^2 \tilde{G}(k_2) - \hat{k}_1^2 \tilde{G}(k_1)}{\tilde{G}(k_1) - \tilde{G}(k_2)}, \quad \hat{k}^2 = \sum_\mu 4 \sin^2(k_\mu/2), \quad (13)$$

with $k_1 = (2\pi/L)(1, 0)$ (averaged with $(0, 1)$) and $k_2 = (2\pi/L)(1, 1)$, validated against the exact free-field propagator. The susceptibility χ itself is always the $\langle |M| \rangle$ -subtracted form defined above.

B. Monte Carlo Data Generation

Field configurations are generated using the Metropolis–Hastings algorithm with single-site Gaussian proposals. For lattices of size $L \geq 32$, we employ a *checkerboard decomposition*: sites are partitioned into two sublattices by coordinate parity, and all sites on the same sublattice are updated simultaneously in a single vectorized operation. This preserves detailed balance (no two updated sites are neighbors) while achieving 10–50 $\times$ speedup over sequential updates.

We generate 5000 configurations at each lattice size ($L = 8, 16, 32, 64$) for training, with $m^2 = -0.5$ and

TABLE I. Action prediction accuracy across lattice sizes. Pearson correlation r and mean relative error are computed on held-out validation sets; entries are mean $\pm$ std over 5 seeds.

Lattice	Sites	Pearson r	Rel. Error
8×8	64	0.99999 ± 10^{-6}	$0.05 \pm 0.00\%$
16×16	256	0.99998 ± 10^{-6}	$0.05 \pm 0.00\%$
32×32	1024	0.99988 ± 0.00002	$0.04 \pm 0.00\%$
64×64	4096	0.98142 ± 0.03558	$0.23 \pm 0.32\%$

$\lambda = 0.5$; one model is trained per lattice size in Table I (Sec. IV F tests a single model across sizes). Configurations are separated by $n_{\text{sweeps}} = 10$ full sweeps to reduce autocorrelation, following 1000 thermalization sweeps. Acceptance rates are maintained near the optimal $\sim 47\%$ by tuning the proposal step size.

C. Training

The GNN is trained to minimize the mean squared error between predicted and exact Euclidean actions:

$$\mathcal{L} = \frac{1}{N_{\text{batch}}} \sum_{i=1}^{N_{\text{batch}}} \left(\hat{S}_E[\phi_i] - S_E[\phi_i] \right)^2. \quad (14)$$

Model hyperparameters: hidden dimension $H = 64$, $B = 3$ message-passing blocks, GELU activation, 2-layer MLPs in encoders; 150,337 parameters at 16×16 (150,529 for the coupling-conditioned variant of Sec. IV C). Training uses AdamW (learning rate 10^{-3} , weight decay 10^{-5}) with cosine annealing over 150 epochs, batch size 32, and an 80/20 train/validation split; seeds and all other settings are recorded in the committed run configurations (`configs/paper/`) and per-run provenance files (`results/`). Monte Carlo data generation runs on laptop CPU; training runs on a single Colab GPU (T4-class, ~ 12 min per seed at 16×16 , ~ 63 min at 64×64).

IV. RESULTS

A. Action Prediction

Table I summarizes the action prediction results. The model achieves near-perfect correlation across lattice sizes, demonstrating that the heterogeneous bipartite architecture can faithfully represent the action functional of scalar ϕ^4 theory. At 64×64 , four of five seeds reach $r \geq 0.9989$, while one converges slowly, reaching only $r = 0.910$ within the fixed budget but $r = 0.9998$ when the epoch budget (and hence the cosine annealing schedule) is doubled — slow convergence, not a failure mode. We report the unfiltered 150-epoch statistics, which inflate the quoted standard deviation; the same seed protocol at $L \leq 32$ shows seed-to-seed spreads of 2×10^{-5}

TABLE II. Architecture comparison on 16×16 lattice ($m^2 = -0.5$, $\lambda = 0.5$). All models trained for 150 epochs on the same data split; entries are mean $\pm$ std over 5 seeds.

Model	Params	Pearson r	Rel. Error
HeteroGNN (ours)	150,337	0.99998 ± 10^{-6}	$0.05 \pm 0.00\%$
Homogeneous GNN	67,073	0.99986 ± 0.00003	$0.11 \pm 0.01\%$
Lattice CNN	29,153	0.99501 ± 0.00276	$0.67 \pm 0.25\%$
Lattice CNN (matched)	146,233	0.99953 ± 0.00022	$0.22 \pm 0.05\%$
MLP	197,633	0.34227 ± 0.00934	$7.37 \pm 0.03\%$

or less. The mild decrease in typical accuracy with lattice size is expected: the graph has $16 \times$ more nodes at 64×64 , and the receptive field of $B = 3$ message-passing blocks covers a smaller fraction of the lattice.

B. Architecture Comparison

To isolate the contribution of the bipartite architecture, we compare against three baselines trained on the same 16×16 data with matched training protocols (150 epochs, AdamW, cosine annealing):

- **Homogeneous GNN:** Field values concatenated with spacetime coordinates on a single node type; same message-passing depth ($B = 3$) and hidden dimension ($H = 64$), with a skip connection from the initial encoding to the readout to prevent over-smoothing of the field signal.
- **Lattice CNN:** Periodic-padded 2D convolutions treating $\phi(x)$ as a single-channel image (4 layers, 32 channels), plus a *parameter-matched* variant (4 layers, 72 channels, 146k parameters $\approx$ the HeteroGNN’s 150k) to control for model capacity.
- **MLP:** Flattened field configuration fed through a 3-layer feedforward network ($H = 256$).

Table II shows the results. The most significant finding is a *structural failure mode* in the homogeneous GNN: without an explicit skip connection from its initial encoding to the readout head, message passing progressively destroys the field signal. Figure 2 quantifies this with a depth ablation (three seeds per point): the no-skip homogeneous model is intact at $B = 1$ ($r = 0.9995$), becomes bimodal across seeds at $B = 2-3$ ($r = 0.43-0.46 \pm 0.4$ — individual seeds either partially retain or lose the signal), and collapses completely by $B = 4-6$ ($r = 0.07 \pm 0.09$ and 0.009 ± 0.013). The skip connection is not physically motivated—it is an engineering patch required because field values and geometry features share a single node representation that is progressively smoothed during message passing. The bipartite architecture avoids this failure mode by construction: field nodes maintain a separate representation throughout the network, and Fig. 2 shows both the HeteroGNN and the skip-equipped




FIG. 2. Depth ablation at 16×16 (mean $\pm$ std over 3 seeds). **Left:** Pearson r vs. number of message-passing blocks B . The homogeneous GNN without a readout skip connection degrades monotonically with depth and collapses below $r = 0.05$ (dotted line) by $B = 6$; the large error bars at $B = 2-3$ reflect bimodal seed-to-seed outcomes. **Right:** $1 - r$ on a log scale, resolving the persistent gap between the HeteroGNN and the skip-equipped homogeneous baseline at every depth.

homogeneous model holding $r > 0.999$ at every depth tested. Note that the per-stage residual connections in the HeteroGNN (Sec. II B) serve a different purpose: they stabilize gradient flow within each message-passing block, whereas the homogeneous skip connection bypasses message passing entirely to rescue a signal that would otherwise be destroyed.

Beyond architectural stability, all spatially structured models achieve $r > 0.99$, confirming that the ϕ^4 action is learnable as a local functional. The HeteroGNN achieves the lowest relative error (0.05%), roughly half that of the homogeneous GNN with its skip connection (0.11%). Matching the CNN’s parameter count to the HeteroGNN’s (146k vs. 150k) narrows but does not close the gap: the matched CNN reaches $r = 0.9995$ and 0.22% relative error, a $4\times$ larger error than the HeteroGNN at equal capacity, indicating the advantage is architectural rather than a parameter-budget artifact. The MLP baseline fails without any spatial structure ($r = 0.34$), confirming that the lattice topology is essential. Seed-to-seed variation is negligible for the graph models ($\sigma_r \lesssim 3 \times 10^{-5}$) and modest for the CNNs, so the ordering in Table II is stable.

C. Coupling Generalization

To test whether the model learns a transferable action *family* rather than a single action landscape, the coupling constants are made model inputs: field-node features become $[\phi_i, m^2, \lambda]$ and the readout MLP is conditioned on (m^2, λ) (adding 192 parameters). A single model is then trained *jointly* on alternating points of a 13-point m^2 grid spanning $[-2.9, -0.3]$ at $\lambda = 0.5$ — both phases and the critical region — with 3000 configurations per point (thermalization ≥ 2000 sweeps; separations scaled by the measured τ_{int} , up to 72 sweeps near criticality). The held-out couplings interleave the training points (interpolation test), and both grid endpoints are excluded

TABLE III. Generalization across coupling values for a single coupling-conditioned model trained jointly on the $\dagger$ -marked m^2 points and evaluated on 1000 held-out configurations per coupling ($*$ = extrapolation beyond the training grid). Mean $\pm$ std over 5 seeds. Relative error is inflated where the action distribution crosses zero (near $m^2 \approx -2.5$); Pearson r is the meaningful metric there.

m^2	Pearson r	Rel. Error
-2.9^*	0.9972 ± 0.0012	$19.52 \pm 2.65\%$
$-2.6833^\dagger$	1.0000 ± 10^{-7}	$0.07 \pm 0.00\%$
-2.4667	0.9999 ± 10^{-5}	$10.34 \pm 1.06\%$
$-2.25^\dagger$	1.0000 ± 10^{-7}	$5.58 \pm 0.75\%$
-2.0333	0.9999 ± 10^{-6}	$0.46 \pm 0.04\%$
$-1.8167^\dagger$	1.0000 ± 10^{-6}	$0.08 \pm 0.00\%$
-1.6	1.0000 ± 10^{-6}	$0.07 \pm 0.01\%$
$-1.3833^\dagger$	1.0000 ± 10^{-6}	$0.05 \pm 0.00\%$
-1.1667	1.0000 ± 10^{-6}	$0.06 \pm 0.01\%$
$-0.95^\dagger$	1.0000 ± 10^{-6}	$0.04 \pm 0.00\%$
-0.7333	1.0000 ± 10^{-6}	$0.08 \pm 0.03\%$
$-0.5167^\dagger$	1.0000 ± 10^{-6}	$0.04 \pm 0.00\%$
-0.3^*	0.9999 ± 10^{-5}	$0.55 \pm 0.34\%$

from training (extrapolation test).

In an earlier version of this experiment, a model trained at a single coupling with no coupling input was evaluated against target actions computed at each new m^2 ; since no function of ϕ alone can match a family of action functionals indexed by a hidden parameter, that protocol measured a partially unlearnable label shift on top of the distribution shift, and the apparent degradation was severe (e.g. $r = 0.58$ at $m^2 = -2.0$). Table III reports the corrected protocol.

Table III shows that with the couplings as inputs, the single model interpolates essentially perfectly: every held-out interpolation coupling reaches $r \geq 0.9999$, *including* $m^2 = -2.47$ in the critical region where the old protocol appeared to fail. Extrapolation beyond the training grid degrades only at the deep broken-phase endpoint ($m^2 = -2.9$: $r = 0.9972 \pm 0.0012$), while the symmetric-phase endpoint remains at $r = 0.9999$. Because the action is a strictly local functional of the field, its prediction does not suffer near criticality once the label family is identified — long-range physics enters the ensemble *distribution*, not the action’s functional form.

D. Phase Transition and Finite-Size Scaling of the Training Ensembles

This and the following subsection characterize the *Monte Carlo ensembles* used throughout this work: they validate the data-generation pipeline against known Ising-universality results, rather than testing the network itself. We perform a coupling sweep over $m^2 \in [-2.8, -1.5]$ at $\lambda = 0.5$ for $L = 16, 32, 64$ using 2000 configurations per point with warm-starting between successive m^2 values.




FIG. 3. Finite-size scaling results for $\lambda = 0.5$. **Left:** Order parameter $|\phi|$ showing progressive sharpening of the phase transition with increasing L . **Center:** Susceptibility χ with peaks growing as $\chi_{\max} \sim L^{\gamma/\nu}$ and converging toward m_c^2 . **Right:** Correlation length ratio ξ/L with crossing of the $L = 16$ and $L = 32$ curves at $m_c^2 = -2.22 \pm 0.24$. Error bars are binned jackknife errors with bin sizes set by the measured τ_{int} .

Figure 3 shows the three canonical finite-size scaling observables. All error bars are binned jackknife errors with bin sizes set by the measured integrated autocorrelation time (Wolff automatic windowing [16]): τ_{int} is 1–3 stored configurations over most of the scan but rises to $\approx 17, 30$, and 100 at the pseudocritical points of $L = 16, 32, 64$ respectively (configurations separated by 10, 20, 40 Metropolis sweeps), quantifying the critical slowing down of the local update.

The order parameter exhibits the expected S-curve behavior, with the transition sharpening progressively from $L = 16$ to $L = 64$. The susceptibility peaks, located by quadratic fits around the maxima, grow with volume: $\chi_{\max} = 9.1(6), 27.0(1.4), 113(106)$ at $m^2 = -2.094(25), -2.139(4), -2.204(20)$ for $L = 16, 32, 64$. A weighted fit of $\ln \chi_{\max}$ vs $\ln L$ gives

$$\gamma/\nu = 1.57 \pm 0.12, \quad (15)$$

within 1.6σ of the exact two-dimensional Ising value $\gamma/\nu = 1.75$; the $L = 64$ peak height carries a large uncertainty precisely because $\tau_{\text{int}} \approx 100$ there leaves only ~ 20 effectively independent configurations.

The ξ/L ratio shows a crossing of the $L = 16$ and $L = 32$ curves at $m_c^2 = -2.22 \pm 0.24$ (parametric bootstrap over the jackknife errors); at the present statistics the $L = 32/L = 64$ crossing is not resolved within errors. Any residual drift of pairwise crossings with L has two sources: corrections to scaling (the leading irrelevant operator shifts pairwise crossings by $O(L^{-\omega})$, so the $L = 16/32$ crossing need not coincide with the asymptotic m_c^2), and critical slowing down of the local Metropolis algorithm, whose autocorrelation times grow as $\tau \sim \xi^z$ with dynamic exponent $z \approx 2$, degrading the effective statistics near the critical point for the largest lattice.

E. Critical Exponent Extraction

From the ξ/L crossing and data collapse analysis (parametric bootstrap over the binned jackknife errors

TABLE IV. Size transfer: a single model trained at 16×16 († , validation split) evaluated at other lattice sizes without retraining. Mean $\pm$ std over 3 seeds.

Variant	Eval. lattice	Pearson r	Rel. Error
Coordinates	8×8	1.0000 ± 10^{-6}	$0.10 \pm 0.00\%$
	$16 \times 16^\dagger$	1.0000 ± 10^{-6}	$0.05 \pm 0.00\%$
	32×32	1.0000 ± 10^{-6}	$0.03 \pm 0.00\%$
	64×64	1.0000 ± 10^{-6}	$0.02 \pm 0.01\%$
Coordinate-free	8×8	1.0000 ± 10^{-6}	$0.08 \pm 0.01\%$
	$16 \times 16^\dagger$	1.0000 ± 10^{-6}	$0.04 \pm 0.00\%$
	32×32	1.0000 ± 10^{-6}	$0.02 \pm 0.00\%$
	64×64	1.0000 ± 10^{-6}	$0.01 \pm 0.00\%$

of every ξ/L point and the crossing uncertainty), we obtain:

$$\nu = 1.10 \pm 0.43 \quad (\text{exact: } \nu = 1), \quad (16)$$

$$m_c^2 = -2.22 \pm 0.24 \quad (\lambda = 0.5). \quad (17)$$

Dedicated lattice studies of two-dimensional ϕ^4 theory locate the critical line to much higher precision (e.g. Ref. [17]); their coupling normalization differs from ours (continuum-limit ratios $f = \lambda/\mu_c^2$ rather than a fixed bare λ at finite spacing), so we do not attempt a quantitative comparison here.

We emphasize that this serves as a *consistency check* rather than a precision measurement: the central value is compatible with the exact Ising result $\nu = 1$, but the large uncertainty precludes a strong quantitative claim. The dominant source of error is not the GNN architecture but the Monte Carlo sampling itself. The local Metropolis algorithm suffers from critical slowing down ($\tau \sim \xi^z$, $z \approx 2$), which degrades effective statistics precisely where they matter most—near the critical point. A cluster algorithm (Wolff [18] or Swendsen–Wang [19]) would reduce z to near zero and dramatically tighten the ν estimate, but is beyond the scope of this initial architecture demonstration. Our purpose here is to show that the bipartite GNN *captures* the physics of the phase transition, not to compete with dedicated finite-size scaling studies.

F. Size Transfer

The models above are trained one-per-lattice-size. To test whether a single learned action functional transfers across volumes, we train at 16×16 only and evaluate at $L \in \{8, 32, 64\}$ without any retraining, for two spacetime-feature variants: absolute coordinates (as in the rest of this paper) and a coordinate-free variant in which spacetime nodes carry a constant feature and all geometry resides in the displacement edge features (restoring exact translation invariance).

Table IV shows that the learned functional transfers essentially perfectly in both directions — down to 8×8

and up to 64×64 , a $64\times$ change in volume — with $r = 1.0000$ at every size for both variants. Relative error *decreases* with lattice size (from 0.10% at $L = 8$ to 0.01–0.02% at $L = 64$), consistent with an accurate learned per-site action density whose fluctuations average out in the extensive sum. The per-site readout (Eq. (8)) is what makes this possible: the same local functional applies at any volume. Notably, even the absolute-coordinate variant transfers, indicating the model learns to ignore the physically irrelevant coordinate inputs; the coordinate-free variant is nevertheless preferable on principle (exact translation invariance) and attains slightly lower relative error at every size.

V. DISCUSSION

A. Advantages of the Bipartite Architecture

The heterogeneous bipartite graph offers several structural advantages over homogeneous approaches:

Multi-field extensibility. In the homogeneous approach, adding a gauge field $U_\mu(x) \in \text{SU}(N)$ requires concatenating its matrix elements with the scalar field value at each node, breaking the natural group structure. In our architecture, gauge fields become a separate node type with their own encoder, connected to spacetime via typed edges. The three-stage message-passing blocks handle arbitrary numbers of field types without modification.

Physical interpretability. The action readout (Eq. 8) decomposes the total action into per-site contributions $\hat{s}(x)$ that can be directly compared with the lattice action density, enabling local error analysis.

Metric information placement. The lattice spacing a is a metric quantity—it defines the physical distance between neighboring sites. We encode it on adjacency edges as displacement vectors $a\hat{e}_\mu$ rather than as a node feature, reflecting the fact that distances are properties of connections, not of individual points. To validate this choice, we tested four architecture variants (lattice spacing in nodes vs. edges, with and without the explicit a^d volume factor in Eq. 8) across lattice sizes 8×8 through 64×64 at varying lattice spacings $a = 1/L$. Two findings emerged. First, the a^d volume scaling is essential: without it, performance collapses as a decreases (Pearson r drops from > 0.999 to -0.26 at 64×64 with $a = 1/64$), because the MLP readout cannot learn to compensate for a scale factor spanning four orders of magnitude. Second, the choice of whether a resides on nodes or edges has negligible effect on accuracy ($r > 0.999$ for both at all sizes), confirming that the placement is a design choice rather than a performance constraint. We choose edges because the metric naturally lives on links, and this design extends cleanly to gauge theories where link variables $U_\mu(x)$ share the same edge representation.

Over-smoothing mitigation. By separating field and spacetime representations, the architecture avoids

mixing all information into a single node type during message passing. Residual connections and layer normalization further stabilize deep stacking.

Interpretation as a learned effective action. The action readout (Eq. 8) decomposes the total action into per-site contributions, and the finite receptive field of B message-passing blocks limits each contribution to information within Ba lattice spacings. The model therefore learns a *local effective action* operator—analogous to a Wilsonian effective action in which modes above a momentum cutoff $\Lambda \sim 1/(Ba)$ have been integrated out. In this view, the message-passing depth B plays the role of an implicit renormalization scale, controlling the resolution at which the theory is represented. This interpretation explains both the model’s success on the local ϕ^4 action and its degradation near criticality, where the relevant physics involves correlations at scales $\xi \gg Ba$ that lie below the effective cutoff.

B. Limitations

Critical slowing down. The Metropolis sampler’s autocorrelation time diverges at the critical point, limiting our ability to precisely determine ν . This is a sampling limitation, not an architectural one.

Receptive field and locality scale. With $B = 3$ blocks, the model’s receptive field spans $\ell = Ba = 3a$, defining a finite *locality scale* for the learned operator. This is sufficient for the ϕ^4 action, which is a strictly local functional of the field and its nearest-neighbor differences — and indeed action prediction remains accurate even at couplings near criticality once the couplings are model inputs (Sec. IV C): the diverging correlation length changes the ensemble *distribution*, not the action’s local functional form. The locality scale becomes a genuine restriction for observables that are themselves nonlocal — correlators at separations $\gg \ell$, Wilson loops of large area, or spectral quantities — which the network cannot represent regardless of coupling. Increasing B enlarges the locality scale but risks over-smoothing, where distinct node representations collapse toward a common value (measured directly in Fig. 2 for the homogeneous baseline). For theories requiring larger locality scales, possible mitigations include multi-scale pooling layers, skip connections across non-adjacent blocks, or attention mechanisms that allow selective long-range communication.

Lattice symmetries. The current architecture does not explicitly enforce discrete lattice symmetries (translations, C_4 rotations, reflections). Instead, these symmetries are learned implicitly from the data, which is generated with periodic boundary conditions that respect the full symmetry group. This is a deliberate design choice, not a limitation: gauge-equivariant architectures [4, 8] — including, recently, gauge-equivariant message passing on graphs [20] — achieve exact symmetry by constraining the network weights and message

functions, but this rigidity restricts them to the symmetry group built in at design time. Our approach prioritizes *architectural flexibility*—the same model structure handles irregular meshes, non-periodic boundaries, and adaptive lattice refinements without modification. For the multi-theory extensibility that motivates this work (gauge fields, fermions, curved spacetimes), this flexibility is essential. Phase II of this program (Sec. VI) will quantify the trade-off directly by measuring the *learned* gauge-invariance error of the bipartite architecture against exactly equivariant designs such as Ref. [20].

VI. FUTURE DIRECTIONS

The bipartite architecture is designed as a foundation for progressively more complex theories:

Phase II: U(1) gauge theory with fermions. Gauge link variables $U_\mu(x) = e^{iA_\mu(x)} \in \text{U}(1)$ and Dirac spinors $\psi(x) \in \mathbb{C}^2$ are added as new field node types with dedicated encoders. The three-stage message passing naturally handles the tripartite field–gauge–spacetime coupling. This phase will benchmark against the Schwinger model (QED in 1+1D), which has exact solutions for comparison.

Phase III: SU(3) Yang–Mills in 4D. The architecture extends to four-dimensional lattices by simply setting the lattice dimension parameter to 4. Gauge link encoders must respect the SU(3) group structure; we plan to use equivariant neural network layers for this purpose. The target is to reproduce established lattice QCD benchmarks (string tension, glueball spectrum).

Phase IV: Wick rotation bridge. Using the GNN’s learned representation of the Euclidean theory, we aim to construct a bridge to Minkowski spacetime via spectral function reconstruction, connecting imaginary-time lattice data to real-time physics.

Phase V: Native Minkowski dynamics. The ultimate goal is real-time simulation of quantum field dynamics in Minkowski spacetime, potentially using complex Langevin methods guided by the GNN’s learned ac-

tion functional.

VII. CONCLUSION

We have introduced a heterogeneous bipartite graph neural network that separates spacetime geometry from field content—a representation inspired by the geometry–matter distinction fundamental to continuum QFT, and absent from existing lattice QFT architectures. On two-dimensional scalar ϕ^4 theory, the model achieves near-perfect action prediction and transfers across a $64\times$ range of lattice volumes without retraining; the underlying Monte Carlo pipeline reproduces the expected Ising-class finite-size scaling ($\nu = 1.10 \pm 0.43$, $\gamma/\nu = 1.57 \pm 0.12$), validating the ensembles the model is trained on. Comparison against homogeneous GNN, CNN, and MLP baselines demonstrates the viability of the bipartite representation for lattice field theory.

The architecture’s key contribution is not improved accuracy on this benchmark but its *extensibility and physical alignment*: the same graph structure and message-passing protocol accommodate scalar, gauge, and fermion fields through modular encoder additions, while preserving the geometry–field distinction throughout the computation. Metric information (lattice spacing) is encoded on adjacency edges as displacement vectors, placing it where distances physically reside and enabling clean integration with gauge link variables in future phases. This provides a unified framework for progressing from toy models to physically relevant theories, with lattice QCD as the ultimate target.

Code and data are available at <https://github.com/crosis19/qft-graph>.

ACKNOWLEDGMENTS

Substantial portions of the code, analysis, and manuscript were developed in collaboration with Claude (Anthropic), an AI assistant. Pre-submission feedback was obtained from ChatGPT (OpenAI) and Gemini (Google), which informed revisions to the experimental design and presentation.

-
- [1] K. G. Wilson, Confinement of quarks, *Physical Review D* **10**, 2445 (1974).
 - [2] G. Kanwar, M. S. Albergo, D. Boyda, K. Cranmer, D. C. Hackett, S. Racanière, D. J. Rezende, and P. E. Shanahan, Equivariant flow-based sampling for lattice gauge theory, *Physical Review Letters* **125**, 121601 (2020).
 - [3] D. Boyda, G. Kanwar, S. Racanière, D. J. Rezende, M. S. Albergo, K. Cranmer, D. C. Hackett, and P. E. Shanahan, Sampling using SU(N) gauge equivariant flows, *Physical Review D* **103**, 074504 (2021).
 - [4] M. Favoni, A. Ipp, D. I. Müller, and D. Schuh, Lattice gauge equivariant convolutional neural networks, *Physical Review Letters* **128**, 032003 (2022).
 - [5] D. Bachtis, G. Aarts, and B. Lucini, Quantum field-theoretic machine learning, *Physical Review D* **104**, 075021 (2021).
 - [6] K. Cranmer, G. Kanwar, S. Racanière, D. J. Rezende, and P. E. Shanahan, Advances in machine-learning-based sampling motivated by lattice quantum chromodynamics, *Nature Reviews Physics* **5**, 526 (2023).
 - [7] M. S. Albergo, G. Kanwar, and P. E. Shanahan, Flow-based generative models for Markov chain Monte Carlo

- in lattice field theory, *Physical Review D* **100**, 034515 (2019).
- [8] D. Luo, S. Chen, G. Carleo, and B. K. Clark, Gauge equivariant neural networks for quantum lattice gauge theories, *Physical Review Letters* **127**, 276402 (2021).
 - [9] A. Apte, C. Córdova, T.-C. Huang, and A. Ashmore, Deep learning lattice gauge theories, *Phys. Rev. B* **110**, 165133 (2024).
 - [10] M. Cranmer, A. Sanchez-Gonzalez, P. Battaglia, R. Xu, K. Cranmer, D. Spergel, and S. Ho, Discovering symbolic and physical laws using graph networks, *arXiv preprint arXiv:2003.04630* (2020).
 - [11] P. W. Battaglia, J. B. Hamrick, V. Bapst, A. Sanchez-Gonzalez, V. Zambaldi, *et al.*, Relational inductive biases, deep learning, and graph networks, in *arXiv preprint arXiv:1806.01261* (2018).
 - [12] M. Schlichtkrull, T. N. Kipf, P. Bloem, R. van den Berg, I. Titov, and M. Welling, Modeling relational data with graph convolutional networks, in *The Semantic Web (ESWC 2018)* (2018) pp. 593–607.
 - [13] Z. Hu, Y. Dong, K. Wang, and Y. Sun, Heterogeneous graph transformer, in *Proceedings of The Web Conference 2020 (WWW)* (2020) pp. 2704–2710.
 - [14] M. Fey and J. E. Lenssen, Fast graph representation learning with PyTorch Geometric, in *ICLR Workshop on Representation Learning on Graphs and Manifolds* (2019).
 - [15] P. Di Francesco, P. Mathieu, and D. Sénéchal, *Conformal Field Theory* (Springer, 1997).
 - [16] U. Wolff, Monte carlo errors with less errors, *Comput. Phys. Commun.* **156**, 143 (2004).
 - [17] D. Schaich and W. Loinaz, Improved lattice measurement of the critical coupling in ϕ_2^4 theory, *Phys. Rev. D* **79**, 056008 (2009).
 - [18] U. Wolff, Collective Monte Carlo updating for spin systems, *Physical Review Letters* **62**, 361 (1989).
 - [19] R. H. Swendsen and J.-S. Wang, Nonuniversal critical dynamics in Monte Carlo simulations, *Physical Review Letters* **58**, 86 (1987).
 - [20] A. Rayat, Y. Li, and G.-W. Chern, Gauge-equivariant graph neural networks for lattice gauge theories, *arXiv preprint* (2026), *arXiv:2604.20797*.